



(<https://colab.research.google.com/github/joanby/python-ml-course/blob/master/notebooks/T4%20-%20%20-%20Linear%20Regression%20-%20Regresión%20lineal%20con%20statsmodel-Colab.ipynb>)

# Regresión lineal simple en Python

## El paquete statsmodel para regresión lineal

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: data = pd.read_csv("../datasets/ads/Advertising.csv")
```

```
In [3]: data.head()
```

Out [3]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

```
In [4]: import statsmodels.formula.api as smf
```

```
In [5]: lm = smf.ols(formula="Sales~TV", data = data).fit()
```

```
In [6]: lm.params
```

Out [6]: Intercept      7.032594  
TV                    0.047537  
dtype: float64

El modelo lineal predictivo sería  $Sales = 7.032594 + 0.047537 * TV$

```
In [7]: lm.pvalues
```

Out [7]: Intercept      1.406300e-35  
TV                    1.467390e-42  
dtype: float64

```
In [8]: lm.rsquared
```

Out [8]: 0.611875050850071

```
In [9]: lm.rsquared_adj
```

Out [9]: 0.6099148238341623

```
In [10]: lm.summary()
```

Out[10]:

OLS Regression Results

Dep. Variable:	Sales	R-squared:	0.612
Model:	OLS	Adj. R-squared:	0.610
Method:	Least Squares	F-statistic:	312.1
Date:	Thu, 14 Jul 2022	Prob (F-statistic):	1.47e-42
Time:	01:15:21	Log-Likelihood:	-519.05
No. Observations:	200	AIC:	1042.
Df Residuals:	198	BIC:	1049.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	7.0326	0.458	15.360	0.000	6.130	7.935
TV	0.0475	0.003	17.668	0.000	0.042	0.053

Omnibus:	0.531	Durbin-Watson:	1.935
Prob(Omnibus):	0.767	Jarque-Bera (JB):	0.669
Skew:	-0.089	Prob(JB):	0.716
Kurtosis:	2.779	Cond. No.	338.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


```
In [52]: sales_pred = lm.predict(pd.DataFrame(data["TV"]))
sales_pred
data.sales_pred = sales_pred
```

```
In [12]: import matplotlib.pyplot as plt
```

```
In [13]: #!/matplotlib inline
#data.plot(kind = "scatter", x = "TV", y ="Sales")
#plt.plot(pd.DataFrame(data["TV"]), sales_pred, c="red", linewidth = 2)
```

```
In [53]: #corregido
%matplotlib inline
plt.scatter(x = data.TV, y= data.Sales)
plt.plot(data.TV, data.sales_pred, color = 'red',linewidth = 2)
```

Out[53]: [



```
In [15]: data["sales_pred"] = 7.032594 + 0.047537*data["TV"]
```

```
In [16]: data["RSE"] = (data["Sales"]-data["sales_pred"])**2
```

```
In [17]: SSD = sum(data["RSE"])
SSD
```

Out[17]: 2102.5305838896525

```
In [18]: RSE = np.sqrt(SSD/(len(data)-2))
RSE
```

Out[18]: 3.258656369238098

```
In [19]: sales_m = np.mean(data["Sales"])
```

```
In [20]: sales_m
```

Out[20]: 14.022500000000003


```
In [21]: error = RSE/sales_m
```

```
In [22]: error
```

Out[22]: 0.2323876890168014

```
In [23]: plt.hist((data["Sales"]-data["sales_pred"]))

Out[23]: (array([ 4., 10., 13., 17., 40., 42., 33., 16., 17.,  8.]),
          array([-8.3860819 , -6.82624404, -5.26640618, -3.70656832, -2.146730
46,
              -0.5868926 ,  0.97294526,  2.53278312,  4.09262098,  5.652458
84,
              7.2122967 ]),
          <BarContainer object of 10 artists>)
```



# Regresión lineal múltiple en Python

## El paquete statsmodel para regresión múltiple

- Sales ~TV
- Sales ~Newspaper
- Sales ~Radio
- Sales ~TV+Newspaper
- Sales ~TV+Radio
- Sales ~Newspaper+Radio
- Sales ~TV+Newspaper+Radio

```
In [24]: #Añadir el Newspaper al modelo existente
lm2 = smf.ols(formula="Sales~TV+Newspaper", data = data).fit()
```

```
In [25]: lm2.params
```

```
Out[25]: Intercept    5.774948
TV                0.046901
Newspaper         0.044219
dtype: float64
```

```
In [26]: lm2.pvalues
```

```
Out[26]: Intercept    3.145860e-22
TV                5.507584e-44
Newspaper         2.217084e-05
dtype: float64
```

$$\text{Sales} = 5.774948 + 0.046901\text{TV} + 0.044219\text{Newspaper}$$

```
In [27]: lm2.rsquared
```

Out[27]: 0.6458354938293273

```
In [28]: lm2.rsquared_adj
```

Out[28]: 0.6422399150864777

```
In [29]: sales_pred = lm2.predict(data[["TV", "Newspaper"]])
```

```
In [30]: sales_pred
```

Out[30]: 0 19.626901  
1 9.856348  
2 9.646055  
3 15.467318  
4 16.837102  
 ...  
195 8.176802  
196 10.551220  
197 14.359467  
198 22.003458  
199 17.045429  
Length: 200, dtype: float64

```
In [31]: SSD = sum((data["Sales"]-sales_pred)**2)
```

```
In [32]: SSD
```

Out[32]: 1918.5618118968273

```
In [33]: RSE = np.sqrt(SSD/(len(data)-2-1))
```

```
In [34]: RSE
```

Out[34]: 3.120719860252885

```
In [35]: error = RSE / sales_m
```

```
In [36]: error
```

Out[36]: 0.22255089037282116

In [37]:

lm2.summary()

Out[37]: OLS Regression Results

Dep. Variable:	Sales	R-squared:	0.646
Model:	OLS	Adj. R-squared:	0.642
Method:	Least Squares	F-statistic:	179.6
Date:	Thu, 14 Jul 2022	Prob (F-statistic):	3.95e-45
Time:	01:15:26	Log-Likelihood:	-509.89
No. Observations:	200	AIC:	1026.
Df Residuals:	197	BIC:	1036.
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	5.7749	0.525	10.993	0.000	4.739	6.811
TV	0.0469	0.003	18.173	0.000	0.042	0.052
Newspaper	0.0442	0.010	4.346	0.000	0.024	0.064

Omnibus:	0.658	Durbin-Watson:	1.969
Prob(Omnibus):	0.720	Jarque-Bera (JB):	0.415
Skew:	-0.093	Prob(JB):	0.813
Kurtosis:	3.122	Cond. No.	410.

Notes:  
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [38]:

*#Añadir la Radio al modelo existente*  
lm3 = smf.ols(formula="Sales~TV+Radio", data = data).fit()

```
In [39]: lm3.summary()
```

Out[39]:

OLS Regression Results

Dep. Variable:	Sales	R-squared:	0.897
Model:	OLS	Adj. R-squared:	0.896
Method:	Least Squares	F-statistic:	859.6
Date:	Thu, 14 Jul 2022	Prob (F-statistic):	4.83e-98
Time:	01:15:26	Log-Likelihood:	-386.20
No. Observations:	200	AIC:	778.4
Df Residuals:	197	BIC:	788.3
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.9211	0.294	9.919	0.000	2.340	3.502
TV	0.0458	0.001	32.909	0.000	0.043	0.048
Radio	0.1880	0.008	23.382	0.000	0.172	0.204

Omnibus:	60.022	Durbin-Watson:	2.081
Prob(Omnibus):	0.000	Jarque-Bera (JB):	148.679
Skew:	-1.323	Prob(JB):	5.19e-33
Kurtosis:	6.292	Cond. No.	425.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [40]: sales_pred = lm3.predict(data[["TV", "Radio"]])
SSD = sum((data["Sales"]-sales_pred)**2)
RSE = np.sqrt(SSD/(len(data)-2-1))
```

```
In [41]: RSE
```

Out[41]: 1.681360912508001

```
In [42]: RSE/sales_m
```

Out[42]: 0.11990450436855059

```
In [43]: #Añadir la Radio al modelo existente
lm4 = smf.ols(formula="Sales~TV+Radio+Newspaper", data = data).fit()
```

In [44]:

lm4.summary()

Out[44]:

OLS Regression Results

Dep. Variable:	Sales	R-squared:	0.897
Model:	OLS	Adj. R-squared:	0.896
Method:	Least Squares	F-statistic:	570.3
Date:	Thu, 14 Jul 2022	Prob (F-statistic):	1.58e-96
Time:	01:15:26	Log-Likelihood:	-386.18
No. Observations:	200	AIC:	780.4
Df Residuals:	196	BIC:	793.6
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.9389	0.312	9.422	0.000	2.324	3.554
TV	0.0458	0.001	32.809	0.000	0.043	0.049
Radio	0.1885	0.009	21.893	0.000	0.172	0.206
Newspaper	-0.0010	0.006	-0.177	0.860	-0.013	0.011

Omnibus:	60.414	Durbin-Watson:	2.084
Prob(Omnibus):	0.000	Jarque-Bera (JB):	151.241
Skew:	-1.327	Prob(JB):	1.44e-33
Kurtosis:	6.332	Cond. No.	454.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [45]:

sales\_pred = lm4.predict(data[["TV", "Radio", "Newspaper"]])  
SSD = sum((data["Sales"]-sales\_pred)\*\*2)  
RSE = np.sqrt(SSD/(len(data)-3-1))

In [46]:

RSE

Out[46]: 1.6855103734147439

In [47]:

RSE/sales\_m

Out[47]: 0.12020041885646236



# Multicolinealidad

## Factor Inflación de la Varianza

- $VIF = 1$  : Las variables no están correlacionadas
- $VIF < 5$  : Las variables tienen una correlación moderada y se pueden quedar en el modelo
- $VIF > 5$  : Las variables están altamente correlacionadas y deben desaparecer del modelo.

```
In [48]: # Newspaper ~ TV + Radio -> R^2 VIF = 1/(1-R^2)
lm_n = smf.ols(formula="Newspaper~TV+Radio", data = data).fit()
rsquared_n = lm_n.rsquared
VIF = 1/(1-rsquared_n)
VIF
```

Out[48]: 1.1451873787239286

```
In [49]: # TV ~ Newspaper + Radio -> R^2 VIF = 1/(1-R^2)
lm_tv = smf.ols(formula="TV~Newspaper+Radio", data=data).fit()
rsquared_tv = lm_tv.rsquared
VIF = 1/(1-rsquared_tv)
VIF
```

Out[49]: 1.00461078493965

```
In [50]: # Radio ~ TV + Newspaper -> R^2 VIF = 1/(1-R^2)
lm_r = smf.ols(formula="Radio~Newspaper+TV", data=data).fit()
rsquared_r = lm_r.rsquared
VIF = 1/(1-rsquared_r)
VIF
```

Out[50]: 1.1449519171055351

In [51]:

lm3.summary()

Out[51]: OLS Regression Results

Dep. Variable:	Sales	R-squared:	0.897
Model:	OLS	Adj. R-squared:	0.896
Method:	Least Squares	F-statistic:	859.6
Date:	Thu, 14 Jul 2022	Prob (F-statistic):	4.83e-98
Time:	01:15:26	Log-Likelihood:	-386.20
No. Observations:	200	AIC:	778.4
Df Residuals:	197	BIC:	788.3
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.9211	0.294	9.919	0.000	2.340	3.502
TV	0.0458	0.001	32.909	0.000	0.043	0.048
Radio	0.1880	0.008	23.382	0.000	0.172	0.204

Omnibus:	60.022	Durbin-Watson:	2.081
Prob(Omnibus):	0.000	Jarque-Bera (JB):	148.679
Skew:	-1.323	Prob(JB):	5.19e-33
Kurtosis:	6.292	Cond. No.	425.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.